In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model\_selection import train\_test\_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
import torch
import torch.nn as nn
import torch.optim
from torch.utils.data import DataLoader, TensorDataset

**READING IN THE MODEL** 

In [2]:	<pre>df = pd.read_csv(r'/kaggle/input/yeast-dataset/yeast.xls')</pre>
	df

Out[2]:		mcg	gvh	alm	mit	erl	рох	vac	nuc	name
	0	0.58	0.61	0.47	0.13	0.5	0.0	0.48	0.22	MIT
	1	0.43	0.67	0.48	0.27	0.5	0.0	0.53	0.22	MIT
	2	0.64	0.62	0.49	0.15	0.5	0.0	0.53	0.22	MIT
	3	0.58	0.44	0.57	0.13	0.5	0.0	0.54	0.22	NUC
	4	0.42	0.44	0.48	0.54	0.5	0.0	0.48	0.22	MIT
	•••									
	1479	0.81	0.62	0.43	0.17	0.5	0.0	0.53	0.22	ME2
	1480	0.47	0.43	0.61	0.40	0.5	0.0	0.48	0.47	NUC
	1481	0.67	0.57	0.36	0.19	0.5	0.0	0.56	0.22	ME2
	1482	0.43	0.40	0.60	0.16	0.5	0.0	0.53	0.39	NUC
	1483	0.65	0.54	0.54	0.13	0.5	0.0	0.53	0.22	CYT
	1484 ro	ows × S	olur)	nns						

In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1484 entries, 0 to 1483
Data columns (total 9 columns):
 # Column Non-Null Count Dtype

#	Column	Non-Null Count	Dtype
0	mcg	1484 non-null	float64
1	gvh	1484 non-null	float64
2	alm	1484 non-null	float64
3	mit	1484 non-null	float64
4	erl	1484 non-null	float64
5	pox	1484 non-null	float64
6	vac	1484 non-null	float64
7	nuc	1484 non-null	float64
8	name	1484 non-null	object

dtypes: float64(8), object(1)
memory usage: 104.5+ KB

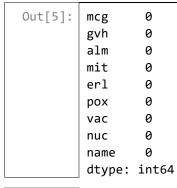
### **DESCRIPTION**

In [4]: df.describe()

Out[4]:		mcg	gvh	alm	mit	erl	рох	vac	nuc
	count	1484.000000	1484.000000	1484.000000	1484.000000	1484.000000	1484.000000	1484.000000	1484.000000
	mean	0.500121	0.499933	0.500034	0.261186	0.504717	0.007500	0.499885	0.276199
	std	0.137299	0.123924	0.086670	0.137098	0.048351	0.075683	0.057797	0.106491
	min	0.110000	0.130000	0.210000	0.000000	0.500000	0.000000	0.000000	0.000000
	25%	0.410000	0.420000	0.460000	0.170000	0.500000	0.000000	0.480000	0.220000
	50%	0.490000	0.490000	0.510000	0.220000	0.500000	0.000000	0.510000	0.220000
	75%	0.580000	0.570000	0.550000	0.320000	0.500000	0.000000	0.530000	0.300000
	max	1.000000	1.000000	1.000000	1.000000	1.000000	0.830000	0.730000	1.000000


# **CHECKING FOR MISSING VALUES**

In [5]: df.isna().sum()



### **HIST PLOT OF FEATURES**

In [6]: | df.hist(figsize=(16, 6), color='blue');



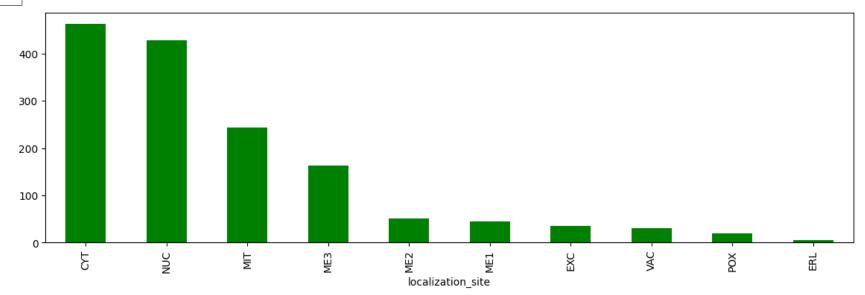
# RENAMING THE TARGET COLUMN

In [7]: df.rename(columns = {'name':'localization\_site'}, inplace=True)

#### TARGET COLUMN DESCRIPTION

```
In [8]: |print(df['localization_site'].value_counts())
        df['localization_site'].value_counts().plot(kind='bar', figsize=(14, 4), color='green')
       localization_site
       CYT
              463
       NUC
              429
       MIT
              244
       ME3
              163
       ME2
               51
       ME1
               44
       EXC
               35
       VAC
               30
       POX
               20
       ERL
       Name: count, dtype: int64
```

Out[8]: <Axes: xlabel='localization\_site'>



### **CORRELATION MATRIX**

```
In [9]: correlation = df.drop('localization_site', axis=1).corr()
plt.figure(figsize=(16, 6))
```



## **SPLITTING TARGET AND INPUT FEATURES**

In [10]: X = df.drop('localization\_site', axis=1) # INPUT FEATURES
y = df['localization\_site'] # TARGET FEATURE

print("Input shape:", X.shape)
print("Target shape:", y.shape)

Input shape: (1484, 8)
Target shape: (1484,)

**ENCODING THE TARGET COLUMN** 

```
In [11]: label_encoder = LabelEncoder()  # LABEL ENCODER

y_encoded = label_encoder.fit_transform(y)  # FITTING THE TARGET COLUMN TO THE ENCODER

classes_names = label_encoder.classes_  # TARGET COLUMN CLASSES
```

#### SCALING THE INPUT FEATURES

```
In [12]: scaler = StandardScaler()  # SCALER
X_scaled = scaler.fit_transform(X) # FITTING THE INPUT FEATURES
```

#### SPLITTING THE TRAIN AND TEST SETS AND CONVERTING TO TENSORS

```
In [13]:
       from sklearn.model_selection import train_test_split
        import torch
       # Split the dataset into training and testing sets
        # -----
       # - X scaled: your normalized or standardized feature matrix
       # - y encoded: your encoded labels (numeric form)
       # - test size=0.2 means 20% test, 80% train
       # - random state ensures reproducibility
       # - stratify=y encoded ensures class balance in both splits
       X_train, X_test, y_train, y_test = train_test_split(
           X scaled,
           y encoded,
           test size=0.2,
           random_state=42,
           stratify=y_encoded
         ______
        # Convert NumPy arrays to PyTorch tensors
        # -----
        # Convert features (X) to float32 tensors — required for neural networks
       X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
        X test tensor = torch.tensor(X test, dtype=torch.float32)
        # Convert labels (y) to long tensors — required by nn.CrossEntropyLoss
       y_train_tensor = torch.tensor(y_train, dtype=torch.long)
       y_test_tensor = torch.tensor(y_test, dtype=torch.long)
```

```
X_train shape: torch.Size([1187, 8])
X_test shape: torch.Size([297, 8])
y_train shape: torch.Size([1187])
y_test shape: torch.Size([297])
Length of classes: 10
X_train shape: torch.Size([1187, 8])
X_test shape: torch.Size([297, 8])
y_train shape: torch.Size([1187])
y_test shape: torch.Size([297])
Length of classes: 10
```

#### **BUILDING A NEURAL NETWORK**

```
In [15]: | # -----
     # Full-Batch Gradient Descent using Adam Optimizer in PyTorch
     import torch
     import torch.nn as nn
     # -----
     # Define loss function and optimizer
     # -----
     model_ad = yeast_model()
     criterion = nn.CrossEntropyLoss()
                                       # Multi-class classification loss
     optimizer = torch.optim.Adam(model_ad.parameters(), lr=0.01) # Adam optimizer with learning rate = 0.1
     # -----
     # Initialize metric storage
     # -----
     train_losses = []
     train accuracies = []
     num_epoch = 100 # Number of training iterations (epochs)
     # -----
     # Training Loop
     # ------
```

```
for epoch in range(num_epoch):
  # -----
  # Forward Pass: Compute predictions and loss
  # -----
  logits = model_ad(X_train_tensor)  # Full-batch forward pass
  loss = criterion(logits, y_train_tensor) # Compute cross-entropy loss
  # -----
  # Backward Pass: Gradient computation and update
  # -----
  optimizer.zero_grad() # Reset accumulated gradients
  loss.backward() # Compute new gradients
              # Update model parameters
  optimizer.step()
  # -----
  # Track metrics
  # -----
  train_losses.append(loss.item())
  # Compute accuracy
  preds = torch.argmax(logits, dim=1)
                                     # Get predicted classes
  correct = (preds == y_train_tensor).sum().item() # Count correct predictions
  accuracy = correct / len(y_train_tensor)
                                       # Compute accuracy
  train_accuracies.append(accuracy)
  # ------
  # Print progress every 10 epochs
  # -----
  if (epoch + 1) % 10 == 0:
     print(f"Epoch {epoch + 1:3d} | Loss: {loss.item():.3f} | Train Acc: {accuracy * 100:.2f}%")
```

```
Epoch 10 | Loss: 1.569 | Train Acc: 47.94%

Epoch 20 | Loss: 1.237 | Train Acc: 55.60%

Epoch 30 | Loss: 1.124 | Train Acc: 58.97%

Epoch 40 | Loss: 1.063 | Train Acc: 60.91%

Epoch 50 | Loss: 1.028 | Train Acc: 62.09%

Epoch 60 | Loss: 1.005 | Train Acc: 62.85%

Epoch 70 | Loss: 0.988 | Train Acc: 63.35%

Epoch 80 | Loss: 0.973 | Train Acc: 63.52%

Epoch 90 | Loss: 0.961 | Train Acc: 63.44%

Epoch 100 | Loss: 0.949 | Train Acc: 64.20%
```

YEAST-LOCALIZATION-SITE

#### **OPTIMIZATION**

```
In [16]: import torch
       import torch.nn as nn
       # -----
       # Define a simple feedforward neural network for the Yeast dataset
       # ------
       class yeast_model(nn.Module):
          def __init__(self, input_dim ,hidden_layer, output_dim):
              # Initialize the parent nn.Module class
              super(yeast_model, self).__init__()
              # Define the model architecture using nn. Sequential
              # - Input Layer: 8 features → 32 hidden neurons
              # - Activation: ReLU (introduces non-linearity)
              # - Output layer: 32 → 10 neurons (for 10 yeast classes)
              self.model = nn.Sequential(
                 nn.Linear(8, 32), # Fully connected layer (input → hidden)
                 nn.ReLU(), # Non-linear activation
                 nn.Linear(32, 10) # Fully connected layer (hidden → output)
          # Define the forward pass (how data flows through the network)
          def forward(self, x):
              return self.model(x) # Pass input through the defined layers
       input_dim = X_train_tensor.shape[1] # 8 features
       hidden_layer = 16
       output_dim = len(label_encoder.classes_)
       # -----
       # Instantiate the model
       # -----
       model = yeast_model(input_dim, hidden_layer, output_dim)
       # Print model architecture summary
       print(model)
```

```
yeast_model(
        (model): Sequential(
          (0): Linear(in_features=8, out_features=32, bias=True)
          (1): ReLU()
          (2): Linear(in_features=32, out_features=10, bias=True)
In [17]: | # Combine features and labels into a single training dataset
        Train dataset = TensorDataset(X train tensor, y train tensor)
        # Combine features and labels into a single test dataset
        Test dataset = TensorDataset(X test tensor, y test tensor)
        # Create DataLoader for training set (batches of 32, shuffled each epoch)
        train loader = DataLoader(Train dataset, batch size=32, shuffle=True)
        # Create DataLoader for test set (batches of 32, no shuffling)
        test loader = DataLoader(Test dataset, batch size=32, shuffle=False)
        # Check shapes and class info
        print(" Train set:", X train tensor.shape)
        print(" Test set:", X test tensor.shape)
       print(" Number of Classes:", len(label encoder.classes ))
        print(" Class Names:", label encoder.classes )
       Train set: torch.Size([1187, 8])
       Test set: torch.Size([297, 8])
       Number of Classes: 10
       Class Names: ['CYT' 'ERL' 'EXC' 'ME1' 'ME2' 'ME3' 'MIT' 'NUC' 'POX' 'VAC']
        ______
        NOTE: Understanding SGD, Full-Batch, and Mini-Batch Optimizers in ML
        _______
        11 Full-Batch Gradient Descent
```

### Description:

- Uses the **entire training dataset** to compute the gradient at each step.
- The model parameters are updated **only once per epoch**.

### Advantages:

- Produces a **stable and accurate** estimate of the gradient.
- Guarantees convergence for convex problems.

### Disadvantages:

- Computationally expensive for large datasets.
- Slow training since it processes all samples before one update.

-----

# **2** Stochastic Gradient Descent (SGD)

-----

#### Description:

- Updates model parameters after computing the gradient on a single sample.
- The model updates **more frequently**, introducing randomness (noise).

### - Advantages:

- Fast updates → quicker learning in early epochs.
- Helps **escape local minima** due to random fluctuations.

### - Disadvantages:

- Noisy loss curve (fluctuates a lot).
- Unstable convergence without proper learning rate scheduling.

-----

Mini-Batch Gradient Descent

-----

### Description:

- A compromise between Full-Batch and SGD.
- Divides the training data into **small batches** (e.g., 32, 64, 128 samples).
- Updates model weights after **each batch**.

### Advantages:

- Faster than Full-Batch but more stable than SGD.
- Enables efficient GPU computation (vectorized operations).
- Commonly used in deep learning.

## Disadvantages:

- Requires tuning of batch size.
- Still slightly noisy, though much smoother than SGD.

-----

# **Summary**

\_\_\_\_\_

- Full-Batch → Stable but Slow
- SGD → Fast but Noisy
- Mini-Batch → Balanced (Best for Deep Learning)
- ✓ Most real-world neural networks use Mini-Batch Gradient Descent with adaptive optimizers like Adam, RMSProp, or SGD with Momentum.

-----

# **FULL BATCH GRADIENT DESCENT OPTMIZER**

```
In [18]: # ------
# Full-Batch Gradient Descent (SGD) Training in PyTorch
# ------
import torch
```

```
import torch.nn as nn
# Define optimizer and loss function
model_gd = yeast_model(input_dim, hidden_layer, output_dim)
optimizer = torch.optim.SGD(model.parameters(), lr=0.1) # Full-batch GD
loss_fn = nn.CrossEntropyLoss()
                                   # Classification loss
# Track metrics
train_losses_gd = []
train_accuracy_gd = []
num_epochs = 100 # Number of full passes through the dataset
# -----
# Training Loop
# ------
for epoch in range(num_epochs):
  model.train() # Set model to training mode (important for dropout/batchnorm)
  # -----
  # Forward pass: compute predictions and loss
  # -----
  logits = model(X train tensor) # All samples at once (full-batch)
  loss = loss_fn(logits, y_train_tensor)
  # -----
  # Backward pass: compute gradients and update weights
  # ------
  optimizer.zero_grad() # Added parentheses
  loss.backward()
  optimizer.step()
  # -----
  # Compute accuracy
  # -----
  preds = torch.argmax(logits, dim=1) # Get predicted classes
  correct = (preds == y_train_tensor).float().mean().item() # Mean accuracy
  # -----
  # Store metrics
  # -----
  train_losses_gd.append(loss.item())
```

```
train_accuracy_gd.append(correct)

# Print progress every 10 epochs
if (epoch + 1) % 10 == 0:
    print(f"[FULL-BATCH] Epoch {epoch + 1:3d} | Loss: {loss.item():.4f} | Accuracy: {correct * 100:.2f}%")

[FULL-BATCH] Epoch 10 | Loss: 2.0515 | Accuracy: 30.67%
[FULL-BATCH] Epoch 20 | Loss: 1.8104 | Accuracy: 35.64%
[FULL-BATCH] Epoch 30 | Loss: 1.6436 | Accuracy: 40.19%
[FULL-BATCH] Epoch 40 | Loss: 1.5242 | Accuracy: 47.35%
[FULL-BATCH] Epoch 50 | Loss: 1.4353 | Accuracy: 51.81%
[FULL-BATCH] Epoch 60 | Loss: 1.3671 | Accuracy: 53.16%
[FULL-BATCH] Epoch 70 | Loss: 1.3138 | Accuracy: 53.50%
[FULL-BATCH] Epoch 90 | Loss: 1.2717 | Accuracy: 54.59%
[FULL-BATCH] Epoch 90 | Loss: 1.2380 | Accuracy: 55.60%
[FULL-BATCH] Epoch 100 | Loss: 1.2106 | Accuracy: 56.70%
```

#### STOCHASTIC GRADIENT DESECNT OPTIMIZER

```
In [19]: # STOCHASTIC GRADIENT DESCENT
      # -----
      # Set seed for reproducibility
     # -----
      torch.manual seed(42)
      import random
      # -----
      # Create a new model instance (same structure as before)
      # -----
      # Assuming yeast model is your neural network class (e.g., YeastMLP)
      model sgd = yeast model(input dim, hidden layer, output dim)
      # Define loss function and optimizer
      criterion = nn.CrossEntropyLoss()
                                        # For multi-class classification
      optimizer = torch.optim.SGD(model sgd.parameters(), lr=0.005) # Stochastic Gradient Descent
      # -----
      # Initialize lists to store loss and accuracy over epochs
      # -----
      train losses sgd = []
      train accuracy sgd = []
```

```
# ------
# Convert training tensors into list of (feature, label) pairs
# -----
# This allows us to iterate sample-by-sample (true SGD)
train_samples = list(zip(X_train_tensor, y_train_tensor))
# (Optional) Track training time
# start_time = time.perf_counter()
# -----
# Training loop - one epoch processes all samples individually
# -----
for epoch in range(100):
   epoch loss = 0.0
   correct = 0
   # Shuffle training samples at the start of each epoch
   torch.random.manual_seed(epoch) # Ensures deterministic shuffling
   shuffled_samples = train_samples.copy()
   random.shuffle(shuffled_samples)
   # Iterate through each sample (stochastic update per sample)
   # -----
   for x_sample, y_sample in shuffled_samples:
      # Add batch dimension since model expects 2D input: (batch_size, features)
      x_sample = x_sample.unsqueeze(0)
      y_sample = y_sample.unsqueeze(0)
      # Forward pass
      outputs = model_sgd(x_sample)
      loss = criterion(outputs, y_sample)
      # Backward pass and parameter update
      optimizer.zero_grad() # Clear previous gradients
      loss.backward() # Compute gradients for this sample
      optimizer.step() # Update weights immediately
      # Track metrics
      epoch_loss += loss.item()
      _, predicted = torch.max(outputs.data, 1)
```

```
[SGD] Epoch 10 | Loss: 1.0293 | Accuracy: 61.16% [SGD] Epoch 20 | Loss: 0.9823 | Accuracy: 62.59% [SGD] Epoch 30 | Loss: 0.9569 | Accuracy: 62.51% [SGD] Epoch 40 | Loss: 0.9468 | Accuracy: 63.77% [SGD] Epoch 50 | Loss: 0.9320 | Accuracy: 64.70% [SGD] Epoch 60 | Loss: 0.9174 | Accuracy: 65.12% [SGD] Epoch 70 | Loss: 0.9025 | Accuracy: 65.12% [SGD] Epoch 80 | Loss: 0.8967 | Accuracy: 65.88% [SGD] Epoch 90 | Loss: 0.8907 | Accuracy: 64.53% [SGD] Epoch 100 | Loss: 0.8850 | Accuracy: 65.04%
```

## MINI BATCH STOCHASTIC GRADIENT DESCENT

```
# Re-initialize model, loss function, and optimizer
# -----
model_mb = yeast_model(input_dim, hidden_layer, output_dim)
criterion = nn.CrossEntropyLoss()
                                            # Suitable for multi-class classification
optimizer_mb = torch.optim.SGD(model_mb.parameters(), lr=learning_rate) # Mini-batch SGD optimizer
# -----
# Create DataLoader for mini-batch processing
# ------
# Combines input features and labels into a TensorDataset
# DataLoader automatically handles batching and shuffling
train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
# -----
# Initialize lists to store loss and accuracy for each epoch
# -----
train_losses_mb = []
train_accuracy_mb = []
# (Optional) Track training time
# start_time = time.perf_counter()
# -----
# Training Loop
# -----
for epoch in range(epochs):
   model_mb.train()  # Set model to training mode
epoch_loss = 0.0  # Accumulate total loss per epoch
correct = 0  # Count correct predictions
   total = 0
                   # Count total samples
   # Iterate through mini-batches
   for X_batch, y_batch in train_loader:
      # Forward pass: compute model output
      outputs = model_mb(X_batch)
      # Compute loss for the batch
      loss = criterion(outputs, y_batch)
      # Backward pass: compute gradients
      optimizer_mb.zero_grad() # Clear existing gradients
```

```
loss.backward()
                                   # Compute new gradients
       optimizer_mb.step()
                                   # Update model weights
        # Track total loss (multiply by batch size to get full sample loss)
       epoch_loss += loss.item() * X_batch.size(0)
       # Track accuracy: compare predicted vs actual labels
        _, predicted = torch.max(outputs, 1)
                                                         # Get predicted class (highest logit)
        correct += (predicted == y_batch).sum().item() # Count correct predictions
       total += y_batch.size(0)
                                                         # Count total samples processed
    # Compute average loss and accuracy for the epoch
    avg_loss = epoch_loss / total
    acc = correct / total
    # Store metrics for later visualization
    train_losses_mb.append(avg_loss)
   train_accuracy_mb.append(acc)
    # Print training progress
    print(f"[Mini-Batch] Epoch {epoch+1:2d}: Loss = {avg_loss:.4f} | Accuracy = {acc*100:.4f}")
# (Optional) Measure total training time
# end_time = time.perf_counter()
# print(f"Training completed in {end_time - start_time:.2f} seconds")
```

[Mini-Batch]	Epoch	1:	Loss =	2.2999	Accuracy = 11.4575
[Mini-Batch]	Epoch	2:	Loss =	2.1785	Accuracy = 29.4861
[Mini-Batch]	Epoch	3:	Loss =	2.0747	Accuracy = 35.3833
[Mini-Batch]	Epoch	4:	Loss =	1.9854	Accuracy = 37.0682
[Mini-Batch]	Epoch	5:	Loss =	1.9074	Accuracy = 39.2586
[Mini-Batch]	Epoch	6:	Loss =	1.8371	Accuracy = 39.7641
[Mini-Batch]	Epoch	7:	Loss =	1.7730	Accuracy = $41.0278$
[Mini-Batch]	Epoch	8:	Loss =	1.7142	Accuracy = $42.1230$
[Mini-Batch]	Epoch	9:	Loss =	1.6600	Accuracy = 44.4819
[Mini-Batch]	Epoch	10:	Loss =	1.6109	Accuracy = 46.0826
[Mini-Batch]	Epoch	11:	Loss =	1.5667	Accuracy = $48.2730$
[Mini-Batch]	Epoch	12:	Loss =	1.5258	Accuracy = $49.5366$
[Mini-Batch]	Epoch	13:	Loss =	1.4881	Accuracy = $50.5476$
[Mini-Batch]	Epoch	14:	Loss =	1.4545	Accuracy = 52.7380
[Mini-Batch]	Epoch	15:	Loss =	1.4228	Accuracy = $52.6537$
[Mini-Batch]	Epoch	16:	Loss =	1.3940	Accuracy = $53.4962$
[Mini-Batch]	Epoch	17:	Loss =	1.3677	Accuracy = 53.6647
[Mini-Batch]	Epoch	18:	Loss =	1.3432	Accuracy = $54.4229$
[Mini-Batch]	Epoch	19:	Loss =	1.3219	Accuracy = $54.5072$
[Mini-Batch]	Epoch	20:	Loss =	1.3023	Accuracy = 54.6757
[Mini-Batch]	Epoch	21:	Loss =	1.2847	Accuracy = 54.9284
[Mini-Batch]	Epoch	22:	Loss =	1.2679	Accuracy = $55.2654$
[Mini-Batch]	Epoch	23:	Loss =	1.2524	Accuracy = 55.9393
[Mini-Batch]	Epoch	24:	Loss =	1.2385	Accuracy = $55.5181$
[Mini-Batch]	Epoch	25:	Loss =	1.2259	Accuracy = 56.7818
[Mini-Batch]	Epoch	26:	Loss =	1.2140	Accuracy = $56.6133$
[Mini-Batch]	Epoch	27:	Loss =	1.2031	Accuracy = $57.7928$
[Mini-Batch]	Epoch	28:	Loss =	1.1932	Accuracy = $58.1297$
[Mini-Batch]	Epoch	29:	Loss =	1.1841	Accuracy = 58.8037
[Mini-Batch]	Epoch	30:	Loss =	1.1756	Accuracy = 59.1407
[Mini-Batch]	Epoch	31:	Loss =	1.1678	Accuracy = 59.0564
[Mini-Batch]	Epoch	32:	Loss =	1.1608	Accuracy = 59.1407
[Mini-Batch]	Epoch		Loss =	1.1536	Accuracy = 59.3934
[Mini-Batch]	Epoch	34:	Loss =	1.1474	Accuracy = 59.6462
[Mini-Batch]	Epoch	35:	Loss =	1.1416	Accuracy = 60.3201
[Mini-Batch]	Epoch			1.1359	Accuracy = 59.5619
[Mini-Batch]	Epoch			1.1308	Accuracy = 60.1516
[Mini-Batch]	Epoch			1.1256	Accuracy = 60.2359
[Mini-Batch]	Epoch			1.1210	Accuracy = 60.3201
[Mini-Batch]	Epoch			1.1167	Accuracy = 59.9832
[Mini-Batch]	Epoch			1.1129	Accuracy = 60.3201
[Mini-Batch]	Epoch	42:	Loss =	1.1097	Accuracy = 60.7414

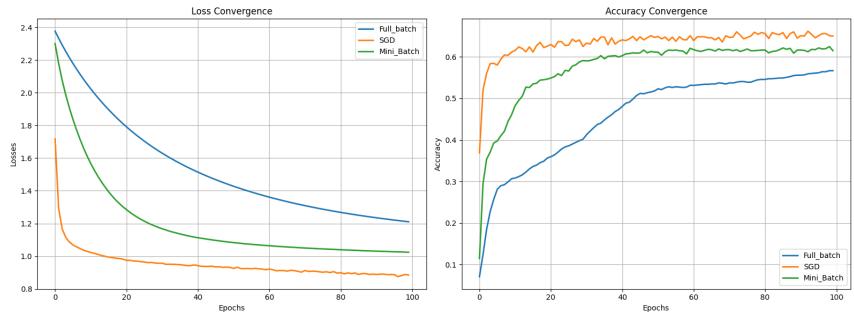
```
[Mini-Batch] Epoch 43: Loss = 1.1063
                                      Accuracy = 60.8256
[Mini-Batch] Epoch 44: Loss = 1.1028 |
                                      Accuracy = 60.9941
[Mini-Batch] Epoch 45: Loss = 1.0997
                                      Accuracy = 60.9099
[Mini-Batch] Epoch 46: Loss = 1.0966
                                      Accuracy = 60.9941
[Mini-Batch] Epoch 47: Loss = 1.0937
                                      Accuracy = 61.6681
[Mini-Batch] Epoch 48: Loss = 1.0908
                                      Accuracy = 60.9941
[Mini-Batch] Epoch 49: Loss = 1.0883
                                      Accuracy = 61.3311
[Mini-Batch] Epoch 50: Loss = 1.0856
                                      Accuracy = 61.1626
[Mini-Batch] Epoch 51: Loss = 1.0834
                                      Accuracy = 61.1626
[Mini-Batch] Epoch 52: Loss = 1.0810
                                      Accuracy = 60.4044
[Mini-Batch] Epoch 53: Loss = 1.0785
                                      Accuracy = 61.2468
[Mini-Batch] Epoch 54: Loss = 1.0764
                                      Accuracy = 61.6681
[Mini-Batch] Epoch 55: Loss = 1.0746
                                      Accuracy = 61.5838
[Mini-Batch] Epoch 56: Loss = 1.0726
                                      Accuracy = 61.6681
[Mini-Batch] Epoch 57: Loss = 1.0707
                                      Accuracy = 61.4996
[Mini-Batch] Epoch 58: Loss = 1.0691
                                      Accuracy = 61.4153
[Mini-Batch] Epoch 59: Loss = 1.0674
                                      Accuracy = 60.7414
[Mini-Batch] Epoch 60: Loss = 1.0654
                                      Accuracy = 62.0893
[Mini-Batch] Epoch 61: Loss = 1.0639
                                      Accuracy = 61.7523
[Mini-Batch] Epoch 62: Loss = 1.0623
                                      Accuracy = 61.4996
[Mini-Batch] Epoch 63: Loss = 1.0610
                                      Accuracy = 61.3311
[Mini-Batch] Epoch 64: Loss = 1.0592
                                      Accuracy = 61.5838
[Mini-Batch] Epoch 65: Loss = 1.0579
                                      Accuracy = 61.8366
[Mini-Batch] Epoch 66: Loss = 1.0565
                                      Accuracy = 61.7523
[Mini-Batch] Epoch 67: Loss = 1.0549
                                      Accuracy = 61.4153
[Mini-Batch] Epoch 68: Loss = 1.0542
                                      Accuracy = 61.9208
[Mini-Batch] Epoch 69: Loss = 1.0525
                                      Accuracy = 61.5838
[Mini-Batch] Epoch 70: Loss = 1.0512
                                      Accuracy = 61.7523
[Mini-Batch] Epoch 71: Loss = 1.0500
                                      Accuracy = 61.7523
[Mini-Batch] Epoch 72: Loss = 1.0488
                                      Accuracy = 61.4996
[Mini-Batch] Epoch 73: Loss = 1.0478
                                      Accuracy = 61.7523
[Mini-Batch] Epoch 74: Loss = 1.0468
                                      Accuracy = 61.3311
[Mini-Batch] Epoch 75: Loss = 1.0455
                                      Accuracy = 61.5838
[Mini-Batch] Epoch 76: Loss = 1.0447
                                      Accuracy = 61.9208
[Mini-Batch] Epoch 77: Loss = 1.0435
                                      Accuracy = 61.4996
[Mini-Batch] Epoch 78: Loss = 1.0424
                                      Accuracy = 61.4996
[Mini-Batch] Epoch 79: Loss = 1.0419
                                      Accuracy = 61.5838
[Mini-Batch] Epoch 80: Loss = 1.0407
                                      Accuracy = 61.6681
[Mini-Batch] Epoch 81: Loss = 1.0397
                                      Accuracy = 61.6681
[Mini-Batch] Epoch 82: Loss = 1.0386
                                      Accuracy = 60.9941
[Mini-Batch] Epoch 83: Loss = 1.0377
                                      Accuracy = 61.3311
[Mini-Batch] Epoch 84: Loss = 1.0368 | Accuracy = 61.4153
```

```
[Mini-Batch] Epoch 85: Loss = 1.0359 | Accuracy = 61.7523
[Mini-Batch] Epoch 86: Loss = 1.0350 | Accuracy = 62.1735
[Mini-Batch] Epoch 87: Loss = 1.0341 | Accuracy = 61.8366
[Mini-Batch] Epoch 88: Loss = 1.0332 | Accuracy = 62.0893
[Mini-Batch] Epoch 89: Loss = 1.0326 | Accuracy = 60.9099
[Mini-Batch] Epoch 90: Loss = 1.0316 | Accuracy = 61.6681
[Mini-Batch] Epoch 91: Loss = 1.0310 | Accuracy = 61.6681
[Mini-Batch] Epoch 92: Loss = 1.0301 | Accuracy = 61.5838
[Mini-Batch] Epoch 93: Loss = 1.0294 | Accuracy = 61.2468
[Mini-Batch] Epoch 94: Loss = 1.0287 | Accuracy = 61.8366
[Mini-Batch] Epoch 95: Loss = 1.0280 | Accuracy = 61.7523
[Mini-Batch] Epoch 96: Loss = 1.0273 | Accuracy = 62.1735
[Mini-Batch] Epoch 97: Loss = 1.0267 | Accuracy = 61.9208
[Mini-Batch] Epoch 98: Loss = 1.0261 | Accuracy = 62.0051
[Mini-Batch] Epoch 99: Loss = 1.0251 | Accuracy = 62.5105
[Mini-Batch] Epoch 100: Loss = 1.0243 | Accuracy = 61.4996
```

```
In [21]: | # -----
      # 🕡 OPTIMIZER PERFORMANCE VISUALIZATION
      # -----
      # This section visualizes the training behavior of three optimizers:
      # - Full-Batch Gradient Descent (GD)
      # - Stochastic Gradient Descent (SGD)
      # - Mini-Batch SGD
      # The Left graph shows how the Loss decreases over epochs (convergence).
      # The right graph shows how model accuracy improves during training.
       # -----
       epochs = len(train losses)
       plt.figure(figsize=(16, 6))
      # -----
       # 

LOSS CONVERGENCE CURVES
      # -----
      plt.subplot(1, 2, 1)
      plt.plot(range(epochs), train losses gd, label='Full batch', lw=2)
      plt.plot(range(epochs), train losses sgd, label='SGD', lw=2)
      plt.plot(range(epochs), train losses mb, label='Mini Batch', lw=2)
      plt.xlabel('Epochs')
      plt.ylabel('Losses')
      plt.title('Loss Convergence')
```

```
plt.legend()
plt.grid(True)
# ACCURACY CONVERGENCE CURVES
# -----
plt.subplot(1, 2, 2)
plt.plot(range(epochs), train_accuracy_gd, label='Full_batch', lw=2)
plt.plot(range(epochs), train_accuracy_sgd, label='SGD', lw=2)
plt.plot(range(epochs), train_accuracy_mb, label='Mini_Batch', lw=2)
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Accuracy Convergence')
plt.legend()
plt.grid(True)
# # LAYOUT & DISPLAY
plt.tight_layout()
plt.show()
```




NOTE: Interpretation of the Optimizer Comparison Graphs
Loss Convergence
• SGD (Per Sample): Shows the fastest initial loss drop, but the curve is noisy and unstable.
• Mini-Batch SGD: Displays a smooth and fast descent, offering a strong balance between speed and stabilit
• Full-Batch GD: Converges steadily but slowly, as its updates are too infrequent for rapid learning.
Accuracy Over Epochs
<ul> <li>SGD: Achieves high accuracy early, but exhibits fluctuations due to noisy updates.</li> <li>Mini-Batch SGD: Provides consistent and steadily improving accuracy with fewer fluctuations.</li> <li>Full-Batch GD: Improves slowly and often reaches lower final accuracy than the other two optimizers.</li> </ul>
Key Insight
Mini-Batch SGD offers the best compromise among all optimization strategies:
Efficient GPU utilization
• 📉 Stable gradient updates
• 4 Faster convergence than Full-Batch GD
Smoother training than per-sample SGD
✓ This is why Mini-Batch training is the standard approach in modern deep learning!
✓ This is why Mini-Batch training is the standard approach in modern deep learning!  ✓ Next Step: Compare these optimizers on the test set using confusion matrices and **final accuracy metrics.**

#### MODEL EVALUATION ON TEST SET

```
|# -----
In [22]:
       # @ MODEL EVALUATION USING CONFUSION MATRIX & ACCURACY
       # -----
       # This section defines a helper function to evaluate trained models
       # using the following metrics and visualizations:
       # - Accuracy Score (via sklearn)
         - Confusion Matrix visualization for class performance
       # Models evaluated:
       # 1. Adam Optimizer
       # 2. Full-Batch Gradient Descent (GD)
       # 3. Stochastic Gradient Descent (SGD - per sample)
       # 4. Mini-Batch SGD
       # -----
       from sklearn.metrics import accuracy score, confusion matrix, ConfusionMatrixDisplay
       # -----
       # @ EVALUATION FUNCTION
       # -----
       def evaluate model(model, X test, y test, class names, title):
          Evaluate a trained PyTorch model on test data and visualize performance.
          - Displays test accuracy.
          - Shows a confusion matrix heatmap.
          # Set model to evaluation mode
          model.eval()
          # Disable gradient tracking for efficiency
          with torch.no grad():
             outputs = model(X test)
             , preds = torch.max(outputs, 1) # Predicted class indices
          # Convert tensors to numpy arrays for sklearn functions
          y_true = y_test.cpu().numpy()
          y_pred = preds.cpu().numpy()
```


```
# A CALCULATE TEST ACCURACY
   # -----
   acc = accuracy_score(y_true, y_pred)
   print(f" ▼ Test Accuracy ({title}): {acc:.4f}")
   # I CONFUSION MATRIX VISUALIZATION
   # -----
   cm = confusion_matrix(y_true, y_pred)
   disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
   # Display matrix
   fig, ax = plt.subplots(figsize=(10, 8))
   disp.plot(ax=ax, cmap="Blues", xticks_rotation=45)
   plt.title(f"Confusion Matrix - {title}")
   plt.show()
            .....
# # EVALUATE ALL MODELS
# -----
evaluate_model(model_ad, X_test_tensor, y_test_tensor, classes_names, "Adam Optimizer")
evaluate_model(model_gd, X_test_tensor, y_test_tensor, classes_names, "Gradient Descent (Full Batch)")
evaluate_model(model_sgd, X_test_tensor, y_test_tensor, classes_names, "Stochastic Gradient Descent (Per Sample)")
evaluate_model(model_mb, X_test_tensor, y_test_tensor, classes_names, "Mini-Batch SGD")
```

✓ Test Accuracy (Adam Optimizer): 0.5825




28<sup>°</sup> 0F 60 10/28/2025, 3:00 PM


▼ Test Accuracy (Gradient Descent (Full Batch)): 0.0875



▼ Test Accuracy (Stochastic Gradient Descent (Per Sample)): 0.5960



▼ Test Accuracy (Mini-Batch SGD): 0.5926



Overfitting occurs when a model learns the training data too well — including its noise, randomness, and outliers — instead of learning the true underlying patterns.

### When this happens:

- **Training accuracy becomes extremely high.**
- X Validation/Test accuracy stops improving or even worsens.
- X The model performs poorly on unseen (new) data.

⚠ Why Does Overfitting Happen?

### Overfitting commonly arises due to:

- A model that is too complex (many neurons, layers, or parameters).
- Insufficient training data relative to model size.
- Lack of regularization (e.g., no Dropout, no L2 penalty).
- Excessive training epochs, allowing the model to memorize data noise.

### ADDING VALIDATION SET TO CHECK FOR MODEL OVERFITTING

```
- Testing (15%)
       from sklearn.model_selection import train_test_split
       # STEP 🚺 - Split into Training and Temporary (Validation + Test)
       # -----
       X_train, X_temp, y_train, y_temp = train_test_split(
          X_scaled,
          y_encoded,
          random_state=42,
          stratify=y_encoded,
          test_size=0.3 # 30% for validation + test
         ______
       # STEP 🗾 - Split Temporary Data into Validation and Test Sets
       # -----
       X_val, X_test, y_val, y_test = train_test_split(
          X_{temp}
          y_temp,
          random_state=42,
          test_size=0.5, # Half of 30% → 15% each
          stratify=y_temp
       # STEP 🗾 - Display Dataset Shapes and Class Info
       # -----
       print(f" Train set: {X_train.shape}")
       print(f" Validation set: {X_val.shape}")
       print(f"  Test set: {X_test.shape}")
       print(f"  Number of Classes: {len(label_encoder.classes_)}")
       print(f" Class Names: {label_encoder.classes_}")
      ✓ Train set: (1038, 8)
       ✓ Validation set: (223, 8)
      Test set: (223, 8)
       Number of Classes: 10
      Class Names: ['CYT' 'ERL' 'EXC' 'ME1' 'ME2' 'ME3' 'MIT' 'NUC' 'POX' 'VAC']
In [24]:
```

```
ODERONAL STATE OF THE PROPERTY OF THE PROPERT
# / Objective:
            Convert training, validation, and test sets into PyTorch tensors
                suitable for model input and training.
import torch
# STEP 1 - Convert Feature Matrices to Float Tensors
X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
X_val_tensor = torch.tensor(X_val, dtype=torch.float32)
X_test_tensor = torch.tensor(X_test, dtype=torch.float32)
 # -----
# STEP 2 - Convert Labels to Long Tensors (for Classification)
# -----
y_train_tensor = torch.tensor(y_train, dtype=torch.long)
y_val_tensor = torch.tensor(y_val, dtype=torch.long)
y_test_tensor = torch.tensor(y_test, dtype=torch.long)
```

### MITIGATING OVERFITTING USING DROPOUT TECHNIQUE

	=========	==========	:=
NOTE: Mitigating Overfitting Using Dropout Technique			
	=======================================		=
What Is Dropout?			

#### Definition:

Dropout is a **regularization technique** used in neural networks to **reduce overfitting** by randomly "dropping out" (i.e., deactivating) a fraction of neurons during training.

#### - Conceptually:

At each training step, a random subset of neurons is ignored — they neither contribute to the forward pass nor participate in backpropagation. This prevents neurons from becoming overly dependent on specific other neurons.

	Why Dropout Helps Prevent Overfitting
	• Forces the network to <b>learn redundant representations</b> , improving robustness.
	Reduces <b>co-adaptation</b> between neurons.
	• Encourages the network to <b>generalize better</b> on unseen data.
	• Acts as a form of <b>model averaging</b> , since each dropout mask represents a different sub-network.
3	How Dropout Works During Training vs. Inference
	<b>During Training:</b> Randomly zeroes out a fraction of activations (e.g., 20–50%).
	• <b>During Inference:</b> Dropout is <b>disabled</b> — all neurons are active, but their outputs are scaled down by the dropout rate maintain balance.
4	Typical Dropout Rates
	Typical Dropout Rates
	Typical Dropout Rates
	Typical Dropout Rates  Input Layer: 0.1 – 0.2
#	Typical Dropout Rates  Input Layer: 0.1 – 0.2  Hidden Layers: 0.3 – 0.5

```
# Define Neural Network with Dropout to Reduce Overfitting
# -----
class ModelDropout(nn.Module):
  def __init__(self, input_dim, hidden_dim, output_dim, dropout_rate=0.5):
     super(ModelDropout, self).__init__()
     # -----
     # Sequential model architecture
     # -----
     self.net = nn.Sequential(
        nn.Linear(input_dim, hidden_dim), # Input → Hidden Layer
       nn.ReLU(),  # Non-linear activation
nn.Dropout(dropout_rate),  # Dropout regularization
       nn.Linear(hidden_dim, output_dim) # Hidden → Output Layer
  # ------
  # Forward Pass
  # -----
  def forward(self, x):
     return self.net(x)
# -----
# Model Initialization Parameters
# -----
input_dim = X_train_tensor.shape[1]
                                     # Number of input features
hidden dim = 32
                                   # Number of hidden neurons
output_dim = len(label_encoder.classes_)
                               # Number of output classes
dropout_rate = 0.5
                                    # Dropout rate (50%)
# -----
# Instantiate Model
# -----
model_dropout_test = ModelDropout(input_dim, hidden_dim, output_dim, dropout_rate)
print(model dropout test)
```

ModelDropout(

(net): Sequential(

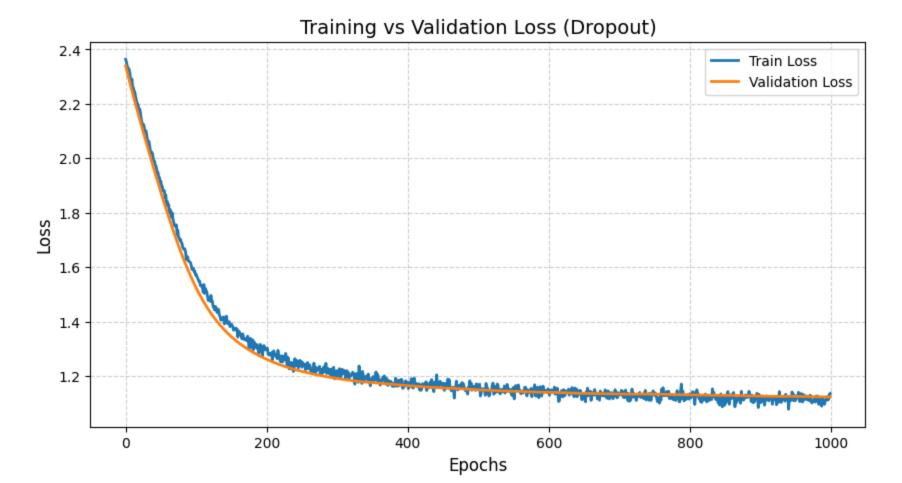
```
(0): Linear(in_features=8, out_features=32, bias=True)
       (1): ReLU()
       (2): Dropout(p=0.5, inplace=False)
       (3): Linear(in_features=32, out_features=10, bias=True)
In [39]:
      # Training MLP Model with Dropout Regularization
      # -----
      import torch
      import torch.nn as nn
      # -----
      # Define Loss Function and Optimizer
      # -----
      criterion = nn.CrossEntropyLoss()
                                                     # Cross-Entropy Loss for multi-class classi
      optimizer = torch.optim.Adam(model dropout test.parameters(), lr=0.001) # Adam optimizer with Learning rate = 0.0
      # -----
      # Initialize Metric Storage
      # -----
      train losses dropout = []
      val losses dropout = []
      epochs = 1000 # Number of training epochs
       _____
      # Training Loop
      # -----
      for epoch in range(epochs):
        # -----
        # Training Phase
        # ------
        model dropout test.train()
                                                 # Enable training mode (activates dropout)
        outputs = model dropout test(X train tensor)
                                                 # Forward pass
        loss = criterion(outputs, y train tensor)
                                                 # Compute training loss
        optimizer.zero_grad()
                                                 # Reset gradients
```

```
loss.backward()
                                              # Backpropagation
optimizer.step()
                                              # Update model parameters
train losses dropout.append(loss.item())
                                                     # Record training loss
# -----
# Validation Phase
# ------
model dropout test.eval()
                                              # Evaluation mode (disables dropout)
with torch.no_grad():
                                              # Disable gradient calculation
   val_outputs = model_dropout_test(X_val_tensor)
                                              # Forward pass on validation data
   val_loss = criterion(val_outputs, y_val_tensor)
                                              # Compute validation loss
   val_losses_dropout.append(val_loss.item())
                                                     # Record validation loss
# -----
# Progress Logging
# ------
if (epoch + 1) % 50 == 0:
   print(f"[Dropout] Epoch {epoch + 1} | Train Loss = {loss.item():.4f} | Val Loss = {val_loss.item():.4f}")
```

[Dropout] Epoch 50 | Train Loss = 1.9119 | Val Loss = 1.8838 [Dropout] Epoch 100 | Train Loss = 1.5720 | Val Loss = 1.5296 [Dropout] Epoch 150 | Train Loss = 1.3886 | Val Loss = 1.3472 [Dropout] Epoch 200 | Train Loss = 1.3015 | Val Loss = 1.2634 [Dropout] Epoch 250 | Train Loss = 1.2547 | Val Loss = 1.2185 [Dropout] Epoch 300 | Train Loss = 1.2204 | Val Loss = 1.1921 [Dropout] Epoch 350 | Train Loss = 1.1930 | Val Loss = 1.1767 [Dropout] Epoch 400 | Train Loss = 1.1782 | Val Loss = 1.1656 [Dropout] Epoch 450 | Train Loss = 1.1641 | Val Loss = 1.1569 [Dropout] Epoch 500 | Train Loss = 1.1605 | Val Loss = 1.1502 [Dropout] Epoch 550 | Train Loss = 1.1544 | Val Loss = 1.1450 [Dropout] Epoch 600 | Train Loss = 1.1540 | Val Loss = 1.1407 [Dropout] Epoch 650 | Train Loss = 1.1293 | Val Loss = 1.1373 [Dropout] Epoch 700 | Train Loss = 1.1300 | Val Loss = 1.1345 [Dropout] Epoch 750 | Train Loss = 1.1320 | Val Loss = 1.1345 [Dropout] Epoch 800 | Train Loss = 1.1160 | Val Loss = 1.1315 [Dropout] Epoch 850 | Train Loss = 1.0931 | Val Loss = 1.1299 [Dropout] Epoch 900 | Train Loss = 1.1111 | Val Loss = 1.1269 [Dropout] Epoch 950 | Train Loss = 1.1099 | Val Loss = 1.1259 [Dropout] Epoch 1000 | Train Loss = 1.1368 | Val Loss = 1.1241

YEAST-LOCALIZATION-SITE

```
-----
import matplotlib.pyplot as plt
# -----
# Plot Loss Curves
# -----
plt.figure(figsize=(10, 5))
                                              # Set figure size
plt.plot(train_losses_dropout, label='Train Loss', linewidth=2)
                                             # Plot training loss
plt.plot(val_losses_dropout, label='Validation Loss', linewidth=2)
                                              # Plot validation loss
# ------
# Graph Formatting
# -----
plt.title("Training vs Validation Loss (Dropout)", fontsize=14)
                                              # Title
plt.xlabel("Epochs", fontsize=12)
                                              # X-axis Label
plt.ylabel("Loss", fontsize=12)
                                              # Y-axis Label
plt.legend()
                                              # Display Legend
plt.grid(True, linestyle='--', alpha=0.6)
                                              # Add gridlines for readability
# -----
# Show Plot
# ------
plt.show()
```



# NOTE: Mitigating Overfitting Using L2 Regularization

-----

**MITIGATING OVERFITTING WITH L2 REGULARIZATION** 

What Is L2 Regularization?

- Definition:

L2 Regularization, also known as <b>weight decay</b> , is a technique to prevent overfitting by <b>penalizing large weights</b> in the neural network.
- Conceptually:
During training, the loss function is modified to include a term proportional to the sum of the squares of all weights:
Loss_total = Loss_original + $\lambda * \Sigma(weights^2)$
where $\lambda$ is the regularization strength.
• Encourages smaller weights, which leads to simpler models.
• Reduces the model's tendency to <b>memorize noise</b> in the training data.
Improves <b>generalization</b> on unseen data.
3 How L2 Works During Training
• During gradient descent, each weight is <b>slightly shrunk</b> in addition to the normal update:
<ul> <li>New weight = weight - learning_rate * (gradient + λ * weight)</li> </ul>
This penalizes large weights while keeping useful patterns.
Typical Usage
Commonly used in linear models, logistic regression, and neural networks.
<ul><li>In PyTorch, you can set L2 regularization with the weight_decay parameter in optimizers:</li></ul>
optimizer = torch.optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-4)

# 5 Key Takeaways

- L2 Regularization is a **simple and effective** method to reduce overfitting.
- Works by **constraining weights**, preventing overly complex models.
- Often combined with **Dropout** for stronger regularization in deep learning.

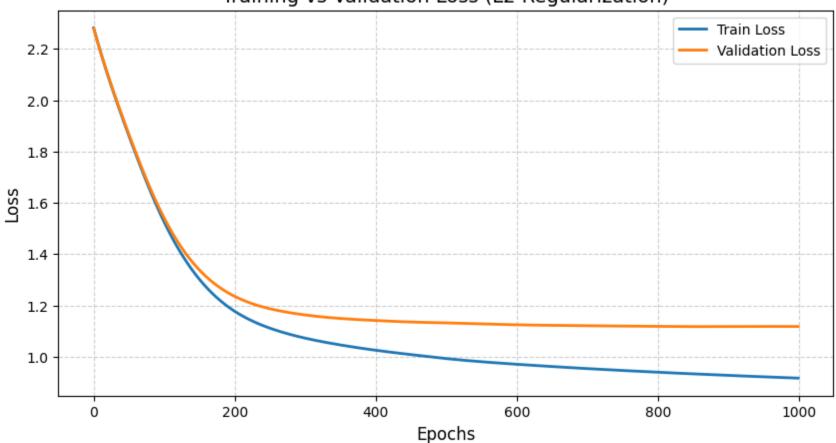
\_\_\_\_\_\_

```
# -----
In [41]:
     # Training MLP Model with L2 Regularization (Weight Decay)
     # -----
     import torch
     import torch.nn as nn
     # -----
     # Define Neural Network with L2 Regularization
     # -----
     class YeastModelL2(nn.Module):
        def __init__(self, input_dim, hidden_dim, output_dim):
          super(YeastModelL2, self).__init__()
          self.network = nn.Sequential(
            nn.Linear(input_dim, hidden_dim), # Input → Hidden Layer
            nn.ReLU(),
                                 # Activation
            nn.Linear(hidden dim, output dim) # Hidden → Output Layer
        # -----
        # Forward Pass
        # -----
        def forward(self, x):
          return self.network(x)
       _____
     # Model, Loss Function, and Optimizer
     # -----
     model 12 = YeastModelL2(
        input dim=X train tensor.shape[1],
```

```
hidden dim=32,
  output_dim=len(label_encoder.classes_)
criterion = nn.CrossEntropyLoss()
                                                    # Multi-class classification loss
optimizer = torch.optim.Adam(model_12.parameters(), lr=0.001, weight_decay=1e-4) # Adam with L2 regularization
# -----
# Initialize Metric Storage
# -----
train losses 12 = []
val losses 12 = []
epochs = 1000
# -----
# Training Loop
# ------
for epoch in range(epochs):
  # -----
  # Training Phase
  # -----
  model 12.train()
                                       # Enable training mode
  outputs_12 = model_12(X_train_tensor)
                                        # Forward pass
  loss_12 = criterion(outputs_12, y_train_tensor) # Compute training loss
  optimizer.zero_grad()
                                       # Reset gradients
  loss 12.backward()
                                        # Backpropagation
  optimizer.step()
                                        # Update model parameters
  train_losses_12.append(loss_12.item())
                                       # Record training loss
  # ------
  # Validation Phase
  model 12.eval()
                                        # Evaluation mode
  with torch.no_grad():
     val_outputs_12 = model_12(X_val_tensor) # Forward pass on validation set
     val_loss_12 = criterion(val_outputs_12, y_val_tensor) # Compute validation loss
     val_losses_12.append(val_loss_12.item()) # Record validation loss
  # -----
  # Progress Logging Every 50 Epochs
```

```
if (epoch + 1) % 50 == 0:
               print(f"[L2 REGULARIZATION] Epoch {epoch + 1} | Train Loss = {loss_l2.item():.4f} | Val Loss = {val_loss_l}
       [L2 REGULARIZATION] Epoch 50 | Train Loss = 1.8688 | Val Loss = 1.8747
       [L2 REGULARIZATION] Epoch 100 | Train Loss = 1.5307 | Val Loss = 1.5463
       [L2 REGULARIZATION] Epoch 150 | Train Loss = 1.3051 | Val Loss = 1.3409
       [L2 REGULARIZATION] Epoch 200
                                    Train Loss = 1.1789 | Val Loss = 1.2367
       [L2 REGULARIZATION] Epoch 250
                                    Train Loss = 1.1122 | Val Loss = 1.1877
       [L2 REGULARIZATION] Epoch 300
                                    Train Loss = 1.0730 | Val Loss = 1.1632
       [L2 REGULARIZATION] Epoch 350
                                    Train Loss = 1.0463 | Val Loss = 1.1494
                                    Train Loss = 1.0256 | Val Loss = 1.1415
       [L2 REGULARIZATION] Epoch 400
       [L2 REGULARIZATION] Epoch 450
                                    Train Loss = 1.0086 | Val Loss = 1.1357
       [L2 REGULARIZATION] Epoch 500
                                    Train Loss = 0.9930 | Val Loss = 1.1320
       [L2 REGULARIZATION] Epoch 550 |
                                    Train Loss = 0.9807 | Val Loss = 1.1280
       [L2 REGULARIZATION] Epoch 600
                                    Train Loss = 0.9708 | Val Loss = 1.1245
                                    Train Loss = 0.9619 | Val Loss = 1.1225
       [L2 REGULARIZATION] Epoch 650
       [L2 REGULARIZATION] Epoch 700
                                    Train Loss = 0.9539 | Val Loss = 1.1209
       [L2 REGULARIZATION] Epoch 750 |
                                    Train Loss = 0.9466 | Val Loss = 1.1195
       [L2 REGULARIZATION] Epoch 800 |
                                    Train Loss = 0.9396 | Val Loss = 1.1185
       [L2 REGULARIZATION] Epoch 850 | Train Loss = 0.9333 | Val Loss = 1.1175
       [L2 REGULARIZATION] Epoch 900 | Train Loss = 0.9273 | Val Loss = 1.1179
       [L2 REGULARIZATION] Epoch 950 | Train Loss = 0.9218 | Val Loss = 1.1181
       [L2 REGULARIZATION] Epoch 1000 | Train Loss = 0.9166 | Val Loss = 1.1179
        # ------
In [42]:
        # Visualization: Training vs Validation Loss (L2 Regularization)
        # ------
        import matplotlib.pyplot as plt
        # Plot Loss Curves
        plt.figure(figsize=(10, 5))
                                                                   # Set figure size
        plt.plot(train_losses_12, label='Train Loss', linewidth=2)
                                                                   # Training loss
        plt.plot(val_losses_12, label='Validation Loss', linewidth=2)
                                                                   # Validation loss
        # ------
        # Graph Formatting
        # -----
        plt.title("Training vs Validation Loss (L2 Regularization)", fontsize=14)
```

### Training vs Validation Loss (L2 Regularization)



MITIGATE OVERFITTING WITH DROPOUT AND L2 REGULARIZATION

### What Is Dropout + L2 Regularization?

- Definition: Dropout combined with L2 regularization is a technique that reduces overfitting by both randomly deactivating neurons during training and penalizing large weights in the network.
- Conceptually: At each training step, some neurons are dropped (Dropout) while all weights are slightly shrunk toward zero (L2), forcing the network to learn robust and simpler representations.

## Why It Helps Prevent Overfitting

- Dropout prevents neurons from co-adapting too strongly, encouraging redundancy in feature learning.
- L2 regularization discourages excessively large weights, simplifying the model.
- Together, they reduce memorization of noise and improve generalization on unseen data.

## 🗾 How It Works During Training vs. Inference

- During Training: Dropout randomly zeroes out activations; L2 adds a penalty proportional to the square of weights.
- During Inference: Dropout is disabled; L2 remains active through the weight constraints applied during training.

## Typical Settings

- Dropout Rate: 0.2 0.5 depending on layer and network size.
- L2 Weight Decay (λ): 1e-4 to 1e-2 is commonly used in neural networks.
- Combined, these hyperparameters may require tuning for optimal performance.

# 5 Key Takeaways

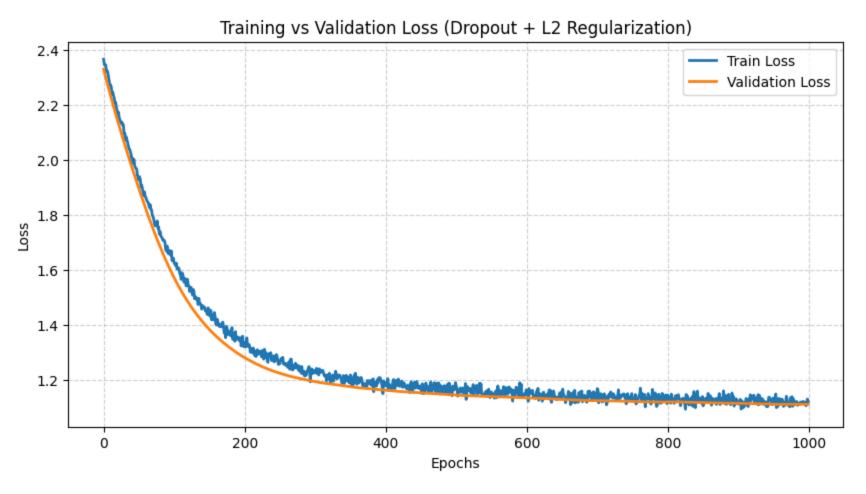
- Dropout + L2 is a strong regularization combo, especially useful for deeper or larger networks.
- It helps the network learn robust features while avoiding overfitting.
- Excessive regularization can cause underfitting, so balance is key.

\_\_\_\_\_\_

```
# -----
In [64]:
     # Model Definition: MLP with Dropout + L2 Regularization
     # -----
     import torch
     import torch.nn as nn
     # -----
     # Define Neural Network with Dropout
     # ------
     class ModelDropoutL2(nn.Module):
        def __init__(self, input_dim, hidden_dim, output_dim, dropout_rate=0.5):
          super(ModelDropoutL2, self).__init__()
          # -----
          # Sequential model architecture
          # -----
          self.net = nn.Sequential(
             nn.Linear(input_dim, hidden_dim), # Input → Hidden Layer
                                 # Non-linear activation
             nn.ReLU(),
             nn.Dropout(dropout_rate),  # Dropout regularization
             nn.Linear(hidden_dim, output_dim) # Hidden → Output Layer
         ______
        # Forward Pass
        # -----
        def forward(self, x):
          return self.net(x)
     # Model Initialization, Loss Function, Optimizer
```

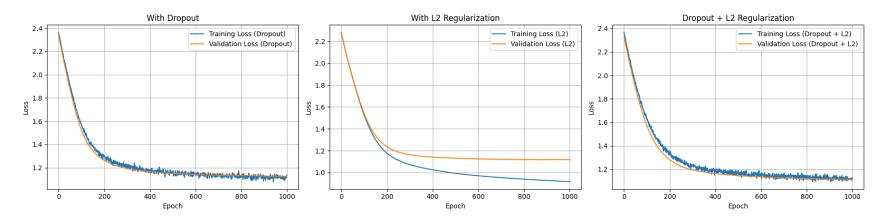
```
input_dim = X_train_tensor.shape[1]
hidden dim = 32
output_dim = len(label_encoder.classes_)
dropout_rate = 0.5
model_dropout_12 = ModelDropoutL2(input_dim, hidden_dim, output_dim, dropout_rate)
criterion = nn.CrossEntropyLoss()
                                                    # Multi-class classification loss
optimizer = torch.optim.Adam(model_dropout_12.parameters(), lr=0.001, weight_decay=1e-4) # L2 regularization via
# -----
# Initialize metric storage
# ------
train_losses_dropout_L2 = []
val_losses_dropout_L2 = []
num_epochs = 1000
# -----
# Training Loop
# -----
for epoch in range(num_epochs):
   # -----
   # Training Step
   # -----
   model_dropout_12.train()
  y_pred = model_dropout_l2(X_train_tensor)  # Forward pass
loss = criterion(y_pred, y_train_tensor)  # Compute loss
   train_losses_dropout_L2.append(loss.item())
                                              # Store train Loss
   optimizer.zero_grad()
                                              # Reset gradients
   loss.backward()
                                             # Backpropagation
   optimizer.step()
                                             # Update weights
   # Validation Step
   # -----
   model_dropout_12.eval()
   with torch.no grad():
      val_output = model_dropout_12(X_val_tensor)
      val_loss = criterion(val_output, y_val_tensor)
      val_losses_dropout_L2.append(val_loss.item())
```

```
______
          # Print progress every 50 epochs
          # -----
          if (epoch + 1) % 50 == 0:
             print(f"[L2+Dropout] Epoch {epoch+1}: Train Loss = {loss.item():.4f}, Val Loss = {val_loss.item():.4f}")
      [L2+Dropout] Epoch 50: Train Loss = 1.8925, Val Loss = 1.8614
      [L2+Dropout] Epoch 100: Train Loss = 1.5747, Val Loss = 1.5452
      [L2+Dropout] Epoch 150: Train Loss = 1.3946, Val Loss = 1.3672
      [L2+Dropout] Epoch 200: Train Loss = 1.2927, Val Loss = 1.2761
      [L2+Dropout] Epoch 250: Train Loss = 1.2314, Val Loss = 1.2298
      [L2+Dropout] Epoch 300: Train Loss = 1.2309, Val Loss = 1.2033
      [L2+Dropout] Epoch 350: Train Loss = 1.1846, Val Loss = 1.1877
      [L2+Dropout] Epoch 400: Train Loss = 1.1788, Val Loss = 1.1771
      [L2+Dropout] Epoch 450: Train Loss = 1.1516, Val Loss = 1.1695
      [L2+Dropout] Epoch 500: Train Loss = 1.1692, Val Loss = 1.1652
      [L2+Dropout] Epoch 550: Train Loss = 1.1625, Val Loss = 1.1604
      [L2+Dropout] Epoch 600: Train Loss = 1.1127, Val Loss = 1.1563
      [L2+Dropout] Epoch 650: Train Loss = 1.1618, Val Loss = 1.1537
      [L2+Dropout] Epoch 700: Train Loss = 1.1266, Val Loss = 1.1516
      [L2+Dropout] Epoch 750: Train Loss = 1.1340, Val Loss = 1.1491
      [L2+Dropout] Epoch 800: Train Loss = 1.1716, Val Loss = 1.1462
      [L2+Dropout] Epoch 850: Train Loss = 1.1033, Val Loss = 1.1448
      [L2+Dropout] Epoch 900: Train Loss = 1.1449, Val Loss = 1.1435
      [L2+Dropout] Epoch 950: Train Loss = 1.1246, Val Loss = 1.1417
      [L2+Dropout] Epoch 1000: Train Loss = 1.1153, Val Loss = 1.1389
       |# -----
In [46]:
       # Plot Training vs Validation Loss for Dropout + L2 Regularization
       # -----
       import matplotlib.pyplot as plt
       # -----
       # Create Figure
       # ------
       plt.figure(figsize=(10, 5))
       # -----
       # Plot Training Loss
       # -----
       plt.plot(train_losses_dropout_L2, label='Train Loss', linewidth=2)
```



```
# -----
axs[0].plot(train losses dropout, label="Training Loss (Dropout)")
axs[0].plot(val losses dropout, label="Validation Loss (Dropout)")
axs[0].set_title("With Dropout")
axs[0].set_xlabel("Epoch")
axs[0].set_ylabel("Loss")
axs[0].legend()
axs[0].grid(True)
# -----
# Plot 2: L2 Regularization
# ------
axs[1].plot(train losses 12, label="Training Loss (L2)")
axs[1].plot(val_losses_12, label="Validation Loss (L2)")
axs[1].set_title("With L2 Regularization")
axs[1].set_xlabel("Epoch")
axs[1].set_ylabel("Loss")
axs[1].legend()
axs[1].grid(True)
# -----
# Plot 3: Dropout + L2 Regularization
# -----
axs[2].plot(train losses dropout L2, label="Training Loss (Dropout + L2)")
axs[2].plot(val_losses_dropout_L2, label="Validation Loss (Dropout + L2)")
axs[2].set_title("Dropout + L2 Regularization")
axs[2].set_xlabel("Epoch")
axs[2].set_ylabel("Loss")
axs[2].legend()
axs[2].grid(True)
# Final Layout Adjustments
# ------
plt.suptitle("Comparison of Regularization Strategies", fontsize=16)
plt.tight_layout(rect=[0, 0, 1, 0.95])
plt.show()
```

#### Comparison of Regularization Strategies



#### **EVALUATION ON TEST SET**

```
In [63]:
         # Evaluate Models on Test Set
         import torch
         from sklearn.metrics import accuracy score, f1 score, precision score
         import pandas as pd
         # Function: Evaluate a Model
         def evaluate_model(model, X_test_tensor, y_test_tensor):
             Evaluates a PyTorch classification model on the test set.
             Returns Accuracy, Macro F1 Score, and Precision.
             model.eval() # Set model to evaluation mode
             with torch.no grad(): # Disable gradient tracking
                 logits = model(X test tensor)
                                                                # Forward pass
                 preds = torch.argmax(logits, dim=1)
                                                               # Predicted class indices
             # Convert to numpy arrays for sklearn metrics
             y_true = y_test_tensor.cpu().numpy()
             y_pred = preds.cpu().numpy()
```

```
# Compute metrics
   acc = accuracy_score(y_true, y_pred)
   f1 = f1_score(y_true, y_pred, average='macro')
   return acc, f1
# -----
# Evaluate Each Regularization Model
# ------
# Dropout Only (M1)
acc_m1, f1_m1, = evaluate_model(model_dropout_test, X_test_tensor, y_test_tensor)
# L2 Regularization Only (M2)
acc_m2, f1_m2 = evaluate_model(model_12, X_test_tensor, y_test_tensor)
# Dropout + L2 (M3)
acc_m3, f1_m3 = evaluate_model(model_dropout_12, X_test_tensor, y_test_tensor)
# Create Results DataFrame
results = {
   "Model": ["Dropout (M1)", "L2 (M2)", "Dropout + L2 (M3)"],
   "Test Accuracy": [acc_m1, acc_m2, acc_m3],
   "Macro F1 Score": [f1_m1, f1_m2, f1_m3],
results_df = pd.DataFrame(results)
# ------
# Display Results
print(results_df)
```

```
    Model
    Test Accuracy
    Macro F1 Score

    0
    Dropout (M1)
    0.573991
    0.451516

    1
    L2 (M2)
    0.605381
    0.605480

    2
    Dropout + L2 (M3)
    0.587444
    0.433184
```

```
In [65]: # ------
```

```
# Save the L2 Regularized Model
# -----
import torch

# Specify the file path where the model will be saved
model_path = "yeast_model_12.pth"

# Save the model's state dictionary
torch.save(model_12.state_dict(), model_path)

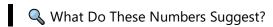
print(f" Model saved successfully at: {model_path}")
```

✓ Model saved successfully at: yeast\_model\_l2.pth

### **Test Performance Summary**

Model	Test Accuracy	Macro F1 Score
Dropout Only (M1)	57.40%	45.15%

| L2 Regularization (M2) | **60.54%** | **60.55%** | Dropout + L2 (M3) | 58.74% | 43.32% |



- ✓ L2 Regularization (M2) performs best across all metrics:
  - Achieved the **highest test accuracy**, **macro F1**, and **macro precision**.
  - Penalizing large weights helped the model generalize better.
- **☑ Dropout Only (M1)** shows moderate performance:
  - Dropout reduced overfitting, but network capacity or training setup may have limited effectiveness.
- ⚠ **Dropout + L2 (M3)** shows slight improvement over Dropout alone in accuracy, but F1 and precision are lower:
  - Combining dropout with L2 didn't yield additive benefits here.
  - Likely due to underfitting from too much regularization.

- Key Takeaways
- **L2 regularization alone (M2)** was most effective for this dataset.
- **Dropout** is more useful in larger networks or datasets prone to severe overfitting.
- Combining Dropout + L2 may require careful tuning of dropout rate and weight decay.
  - What If We Change the Architecture?
- Increase neurons (e.g., 64)?: More expressive model, higher risk of overfitting, requires stronger regularization.
- Decrease neurons (e.g., 16)?: Simpler model, may underfit, training faster but less accurate.
- Add hidden layers?: Can learn deeper representations, needs more data, regularization, and training strategies like batch norm or residual connections.

### Final Reflection: Generalization, Regularization, and Model Design

In this notebook, we investigated how different regularization strategies impact a model's ability to generalize — that is, to perform well on unseen data, not just the training set.

We experimented with Dropout, L2 Weight Decay, and a combination of both. Throughout, we observed how regularization can significantly influence both accuracy and fairness, especially in imbalanced multiclass datasets.

Key observations from our experiments:

- Dropout performed well too, promoting simpler, more stable solutions by randomly deactivating neurons.
- Combining Dropout + L2 did not improve results in this setup likely because the model was overly constrained without careful tuning.
- These results emphasize a central lesson in deep learning: regularization is powerful but not automatic. Its effectiveness depends on the architecture, dataset size, and learning dynamics. A method that works well in one scenario might underperform in another.

- This connects to the concept of the bias-variance trade-off:
- Underfitting occurs when the model is too simple (high bias).
- Overfitting occurs when the model is too complex (high variance).
- The goal is to find the sweet spot a model expressive enough to capture meaningful patterns, yet restrained enough to generalize well.